Document Q&A Test PDF

This PDF exists to test the extract_from_pdf function. It contains two pages, each with a short paragraph of text that pypdf's PdfReader should be able to extract cleanly.

Page Two

Pinecone stores vector embeddings and supports metadata filtering, which is what allows a RAG pipeline to isolate search results to a single user's own documents.